



Machine Learning

LABORATORY: LSTM-RNN In Class

NAME: 黃仕翔

STUDENT ID#: 313513054

Objectives:

- Students will implement a Recurrent Neural Network (RNN) from scratch, without using built-in RNN APIs.
- Students will manually implement:
 - Hidden state updates over a sequence, a simple linear output layer for sentence classification and manual Stochastic Gradient Descent (SGD) updates.
- Students will understand how RNN hidden states summarize past information.
- Students will visualize the loss curve during training and observe model predictions for simple sentences.

Part 1. Instruction

In this assignment, you will implement a basic Vanilla Recurrent Neural Network (RNN) for simple sequence classification using PyTorch, but without using any high-level RNN modules (no *nn.RNN*, no *optim.SGD*, etc.).

You will manually implement:

- A step-by-step hidden state update based on the RNN equation.
- A manual simple output layer to predict whether a sentence is "good" or "bad".
- Manual parameter updates using basic gradient descent.

The general RNN computation is as follows:

$$s_t = \varphi(Ws_{t-1} + Ux_t + b)$$

After the final time step, you apply an output layer:

$$\text{logits} = W_{out}s_T + b_{out}$$

where:

- $W_{out} \in R^{2 \times n}$ = output weight matrix
- $b_{out} \in R^{2 \times 1}$ = output bias
- The prediction is obtained by applying argmax over the logits.

Part 2. Code Template

Step	Procedure
1	<pre>import torch import matplotlib.pyplot as plt # 1. Vocabulary vocab = { "The": 0, "movie": 1, "is": 2, "good": 3, "bad": 4} def word_to_onehot(word): vec = torch.zeros(len(vocab), 1)</pre>



	<pre> vec[vocab[word]] = 1.0 return vec # 2. Training Samples train_data = [("The", "movie", "is", "good"), 1), ("The", "movie", "is", "bad"), 0)] # 3. Define RNN Parameters n = 2 # hidden size m = len(vocab) T = 4 # number of words torch.manual_seed(42) </pre>
2	<pre> # Initialize Weights (n x n), (n x m), (n x 1), (2 x n), (2 x 1) W = None U = None b = None W_out = None b_out = None # 4. Training Setup learning_rate = 0.1 num_epochs = 300 loss_fn = torch.nn.CrossEntropyLoss() loss_history = [] </pre>
3	<pre> # 5. Training Loop for epoch in range(num_epochs): total_loss = 0.0 for sentence, label in train_data: inputs = [word_to_onehot(word) for word in sentence] s_prev = torch.zeros(n, 1) for x_t in inputs: # TODO: Compute s_t s_t = None # TODO s_prev = s_t # Output layer logits = None # TODO logits = logits.view(1, -1) # Loss target = torch.tensor([label]) loss = loss_fn(logits, target) total_loss += loss.item() # Backward loss.backward() </pre>



	<pre> # Manual update (SGD) with torch.no_grad(): # TODO: Update W, U, b, W_out, b_out pass # Zero gradients after updating W.grad.zero_() U.grad.zero_() b.grad.zero_() W_out.grad.zero_() b_out.grad.zero_() loss_history.append(total_loss) if (epoch + 1) % 50 == 0: print(f'Epoch [{epoch+1}/{num_epochs}] - Loss: {total_loss:.4f}') </pre>
4	<pre> # 6. Plot Loss plt.plot(loss_history) plt.xlabel("Epoch") plt.ylabel("Total Loss") plt.title("Training Loss Curve") plt.grid(True) plt.show() </pre>
5	<pre> # 7. Test After Training print("\n==== Testing ====") test_sentences = [["The", "movie", "is", "bad"], ["The", "movie", "is", "good"]] for test_sentence in test_sentences: inputs = [word_to_onehot(word) for word in test_sentence] s_prev = torch.zeros(n, 1) for x_t in inputs: #Forward pass again (same as above) s_t = None # TODO s_prev = s_t logits = None # TODO prediction = torch.argmax(logits) print(f'Sentence: {test_sentence} → Prediction: {prediction.item()}') </pre>

Grading Assignment & Submission (30% Max)

Implementation:

1. (5%) Correctly initialize all weights.
2. (5%) Correctly compute the hidden state update and the output layer (logits).
3. (5%) Correctly update all parameters manually using gradients (manual SGD).
4. (5%) The model correctly predicts "good" = 1 and "bad" = 0 in the test sentences.

Question:



5. (5%) What does the hidden state s_t represent at each time step when processing a sequence of words?
6. (5%) Why are RNNs hard to train?


Submission :

1. Report: Provide your screenshots of your results in the last pages of this PDF File.
2. Code: Submit your complete Python script in either .py or .ipynb format.
3. Upload both your report and code to the E3 system (**Labs7 In Class Assignment**). Name your files correctly:
 - a. Report: StudentID_Lab7_InClass.pdf
 - b. Code: StudentID_Lab7_InClass.py or StudentID_Lab7_InClass.ipynb
4. Deadline: 16:20 PM
5. Plagiarism is **strictly prohibited**. Submitting copied work from other students will result in penalties.

Results and Discussion:

```
Epoch [50/300] - Loss: 0.0816
Epoch [100/300] - Loss: 0.0347
Epoch [150/300] - Loss: 0.0216
Epoch [200/300] - Loss: 0.0156
Epoch [250/300] - Loss: 0.0122
Epoch [300/300] - Loss: 0.0100
```

Training Loss Curve



```
=== Testing ===
Sentence: ['The', 'movie', 'is', 'bad'] -> Prediction: 0
Sentence: ['The', 'movie', 'is', 'good'] -> Prediction: 1
```

5. (5%) What does the hidden state s_t represent at each time step when processing a sequence of words?

at each step, t ,

- $S_t = \tanh(Ws_{t-1} + Ux_t + b)$
- It is a fixed-size vector that summarize all word up to t
- It blend the previous memory S_{t-1} with the current input x_t
- At the final step S_T feed into the output layer to predict "good" and "bad"

6. (5%) Why are RNNs hard to train?

- **Vanishing gradient**: repeated multiplication by w and $\phi'(<1)$ make gradient shrink exponentially, so early input hardly affect the loss.
- **Exploding gradient**: if those factor exceed 1, gradient blow up, causing instability.
- **Sequential dependency**: each step depend on the last, limiting parallelization and make optimization sensitive to hyperparameter.



